





PART 2. DATA MANAGEMENT




www.dilbert.com

scottadams@aol.com



5-7-08 © 2008 Scott Adams, Inc./Dist. by UFS, Inc.





CREATING A DATASET

Topics

- R data structures
 - vectors, matrices, arrays,
data frames, factors, lists
- Data Input
 - text file
 - Excel
 - Stat packages (SAS, SPSS, Stata)
 - DBMS

R Data Structures


(a) scalar




(b) vector




(c) matrix



(d) array




(e) dataframe



Columns can be different modes

(f) list



Vectors

One dimensional arrays

```
a <- c(1, 2, 5, 3, 6, -2, 4)
```

```
b <- c("one", "two", "three")
```

```
c <- c(TRUE, TRUE, TRUE, FALSE, TRUE, FALSE)
```

Vectors(2)

Identifying elements

```
a <- c(1, 2, 5, 3, 6, -2, 4)
```

```
a[3]
```

```
[1] 5
```

```
a[c(1, 3, 5)]
```

```
[1] 1 5 6
```

```
a[2:6]
```

```
[1] 2 5 3 6 -2
```

Matrices

Two dimensional arrays where each element has same mode (numeric, character, or logical)

	[,1]	[,2]	[,3]	[,4]
[1,]	1	6	11	16
[2,]	2	7	12	17
[3,]	3	8	13	18
[4,]	4	9	14	19
[5,]	5	10	15	20

Creating a Matrix

```
y <- matrix(1:20, nrow=5)
```

y

	[,1]	[,2]	[,3]	[,4]
[1,]	1	6	11	16
[2,]	2	7	12	17
[3,]	3	8	13	18
[4,]	4	9	14	19
[5,]	5	10	15	20

ncol=4 would
work too

Creating a Matrix

```
y <- matrix(1:20, nrow=5, byrow=TRUE)
```

y

	[,1]	[,2]	[,3]	[,4]
[1,]	1	2	3	4
[2,]	5	6	7	8
[3,]	9	10	11	12
[4,]	13	14	15	16
[5,]	17	18	19	20

Creating a Matrix

```
data <- c(3615, 365, 2212, 3624, 6315, 4530)
rnames <- c("Alabama", "Alaska", "Arizona")
cnames <- c("Population", "Income")
y <- matrix(data, ncol=2,
             dimnames=list(rnames, cnames))
```

	Population	Income
Alabama	3615	3624
Alaska	365	6315
Arizona	2212	4530

Using Matrix Subscripts

```
x <- matrix(1:10, nrow = 2)
```

```
x
```

	[,1]	[,2]	[,3]	[,4]	[,5]
[1,]	1	3	5	7	9
[2,]	2	4	6	8	10

```
x[2, ]
```

```
[1] 2 4 6 8 10
```

```
x[, 2]
```

```
[1] 3 4
```

```
x[1, 4]
```

```
[1] 7
```

```
x[1, c(4, 5)]
```

```
[1] 7 9
```

Data frame

- Rectangular array of data
- More general than a matrix - different columns can contain different modes of data (numeric, character, etc.)
- Similar to datasets in SAS, SPSS, and Stata

```
mydata <- data.frame( col1, col2, ..., coln )
```

Creating a data frame

```
ptID      <- c(111, 208, 113, 408)
age        <- c(25, 34, 28, 52)
diabetes   <- c("Type1", "Type2", "Type1", "Type1")
status     <- c("Poor", "Improved", "Excellent", "Poor")
ptdata     <- data.frame(ptID, age, diabetes, status)
ptdata
```

	ptID	age	diabetes	status
1	111	25	Type1	Poor
2	208	34	Type2	Improved
3	113	28	Type1	Excellent
4	408	52	Type1	Poor

Specifying elements of a data frame

```
ptdata[1:2]
```

	ptID	age
1	111	25
2	208	34
3	113	28
4	408	52

```
ptdata[c("diabetes","status")]
```

	diabetes	status
1	Type1	Poor
2	Type2	Improved
3	Type1	Excellent
4	Type1	Poor

```
ptdata$age
```

```
[1] 25 34 28 52
```

Specifying elements of a data frame

ptdata[1:2]

	ptID	age
1	111	25
2	208	34
3	113	28
4	408	52

ptdata[c(1,3), 1:2]

	ptID	age
1	111	25
3	113	28

ptdata[2:3, 1:2]

	ptID	age
2	208	34
3	113	28

With

```
summary(mtcars$mpg)  
plot(mtcars$mpg, mtcars$disp)
```

```
with(mtcars, {  
    summary(mpg)  
    plot(mpg, disp)  
})
```

Factors

- Data structure specifying categorical (nominal) or ordered categorical (ordinal) variables
- Tells R how to handle that variable in analyses
- Very important and misunderstood
- Any variable that is categorical or ordinal should usually be stored as a factor.

Factors (2)

```
ptdata$sex <- c(1, 1, 2, 5)
```

```
ptdata$sex <- factor(sex, levels=c(1, 2),  
  labels=c("Male", "Female"))
```

associates 1=Male, 2=Female

Treats sex as a categorical variable in all analyses

What happens to sex=5?

Factors (2)

```
ptdata$status <- c(1, 2, 3, 1)
```

```
ptdata$status <- factor(ptdata$status,  
  ordered=TRUE,  
  levels=c(1, 2, 3)  
  levels=c("Poor", "Improved", "Excellent"))
```

associates 1=Poor, 2=Improved, 3=Excellent

Treats status as an ordinal variable in all analyses

Lists

- Ordered collection of objects (components)
- Many important functions return lists

```
mylist <- list(name1=object1, name2=object2, ...)
```

List Example

mylist

```
g <- "My First List"
h <- c(25, 26, 18, 39)
j <- matrix(1:10, nrow = 5)
k <- c("one", "two", "three")
```

```
mylist <- list(title = g,
               ages = h,
               mymatrix = j,
               mystrings = k)
```

```
$title
[1] "My First List"
```

```
$ages
[1] 25 26 18 39
```

```
$mymatrix
      [,1] [,2]
[1,]   1   6
[2,]   2   7
[3,]   3   8
[4,]   4   9
[5,]   5  10
```

```
$mystrings
[1] "one" "two" "three"
```

List Example

```
mylist[[2]]  
[1] 25 26 18 39
```

```
mylist[["ages"]]  
[1] 25 26 18 39
```

```
mylist$ages  
[1] 25 26 18 39
```

mylist

```
$title  
[1] "My First List"
```

```
$ages  
[1] 25 26 18 39
```

```
$mymatrix  
[,1] [,2]
```

```
[1,] 1 6  
[2,] 2 7  
[3,] 3 8  
[4,] 4 9  
[5,] 5 10
```

```
$mystrings  
[1] "one" "two" "three"
```

List Example

```
mylist$ages[2]
```

```
[1] 26
```

```
mylist[[2]][2]
```

```
[1] 26
```

```
mylist$mymatrix[2,2]
```

```
[1] 7
```

```
mylist[[3]][,2]
```

```
[1] 6 7 8 9 10
```

mylist

```
$title
```

```
[1] "My First List"
```

```
$ages
```

```
[1] 25 26 18 39
```

```
$mymatrix
```

```
  [,1] [,2]
```

```
[1,]  1  6
```

```
[2,]  2  7
```

```
[3,]  3  8
```

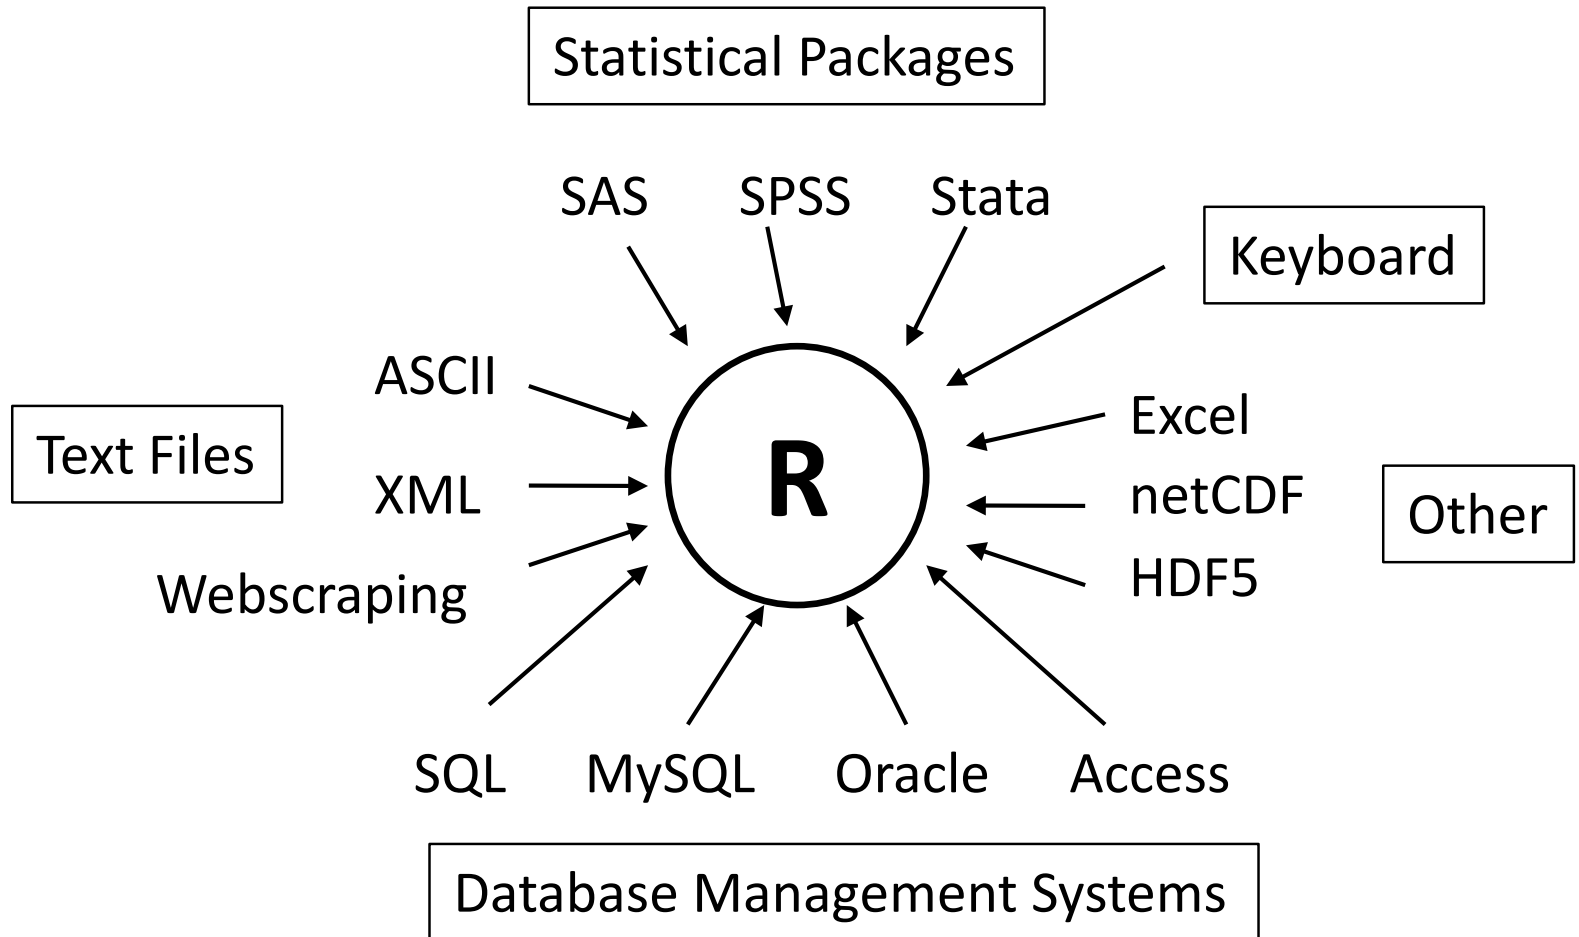
```
[4,]  4  9
```

```
[5,]  5 10
```

```
$mystrings
```

```
[1] "one" "two" "three"
```


Data Input



Import from Delimited Text File

First row are variable names

```
library(readr)
```

```
# comma separated values
```

```
mydataframe <- read_csv("file")
```

```
# tab separated values
```

```
mydataframe <- read_tsv("file")
```

```
# semicolon separated values
```

```
mydataframe <- read_csv2("file")
```

option:
col_names=FALSE

Import from Delimited Text File

First row are variable names

```
library(readr)
```

```
# space delimited values
```

```
mydataframe <- read_table("file")
```

```
# Like read.table(), it allows any number
```

```
# whitespace characters between columns, and
```

```
# the lines can be of different lengths
```

```
mydataframe <- read_table2("file")
```

option:
col_names=FALSE

Import from Delimited Text File

```
library(readr)
```

```
# more control over reading file
```

```
mydataframe <- read_delim("file",  
  delim,  
  col_names = TRUE,  
  na = c("", "NA"),  
  skip = 0,  
  n_max = Inf)
```

Importing from Excel

```
# install.packages(readxl)  
library(readxl)  
df <- read_excel("myfile.xlsx", 1)
```

Importing from Stat Packages

```
# install.packages(haven)
library(haven)

# sas
df <- read_sas("myfile.sas7bdat")

# spss
df <- read_sav("myfile.sav")

# stata
df <- read_stata("myfile.dta")
```

Accessing DBMS

- R can access MS SQL Server, MS Access, MySQL, Oracle, PostgreSQL, DB2, Sybase, Teradata, SQLite, ...
- One of the best ways to deal with large datasets

Functions for working with objects

Function

Action

length(object)

number of elements/components

dim(object)

dimensions of an object

str(object)

structure of an object

class(object)

class or type of an object

names(object)

names of components in an object

c(object, object,...)

combines objects into a vector

cbind(object, object, ...)

combines objects as columns

rbind(object, object, ...)

combines objects as rows

object

prints the object

head(object)

list the first part of the object

tail(object)

list the last part of the object

rm(object)

delete an object


```

> class(mtcars)
[1] "data.frame"
> names(mtcars)
[1] "mpg"  "cyl"  "disp" "hp"    "drat" "wt"    "qsec" "vs"    "am"    "gear"
[11] "carb"
> length(mtcars)
[1] 11
> dim(mtcars)
[1] 32 11
> str(mtcars)
'data.frame':   32 obs. of  11 variables:
 $ mpg : num  21 21 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 ...
 $ cyl : num  6 6 4 6 8 6 8 4 4 6 ...
 $ disp: num  160 160 108 258 360 ...
 $ hp  : num  110 110 93 110 175 105 245 62 95 123 ...
 $ drat: num  3.9 3.9 3.85 3.08 3.15 2.76 3.21 3.69 3.92 3.92 ...
 $ wt  : num  2.62 2.88 2.32 3.21 3.44 ...
 $ qsec: num  16.5 17 18.6 19.4 17 ...
 $ vs  : num  0 0 1 1 0 1 0 1 1 1 ...
 $ am  : num  1 1 1 0 0 0 0 0 0 0 ...
 $ gear: num  4 4 4 3 3 3 3 4 4 4 ...
 $ carb: num  4 4 1 1 2 1 4 2 2 4 ...
> head(mtcars)
      mpg cyl disp  hp drat   wt  qsec vs am gear carb
Mazda RX4           21.0   6  160 110 3.90 2.620 16.46 0  1    4    4
Mazda RX4 Wag       21.0   6  160 110 3.90 2.875 17.02 0  1    4    4
Datsun 710          22.8   4  108  93 3.85 2.320 18.61 1  1    4    1
Hornet 4 Drive      21.4   6  258 110 3.08 3.215 19.44 1  0    3    1
Hornet Sportabout  18.7   8  360 175 3.15 3.440 17.02 0  0    3    2
Valiant             18.1   6  225 105 2.76 3.460 20.22 1  0    3    1

```

Note for programmers

- **The period (.) has no special significance in object names.**

However, the dollar sign (\$) has a somewhat analogous meaning, identifying the parts of an object.

For example, A\$x refers to variable x in data frame A.

- **R does not provide for multi-line or block comments.**

You must start each line of a multi-line comment with #. For debugging purposes, you can also surround code that you want the interpreter to ignore with the statement `if(0){...}`.

- **R does not have scalar values.**

Scalars are represented as one element vectors.

Note for programmers

- **Variables cannot be declared.**
They come into existence on first assignment.
- **Assigning a value to a non-existent element of a vector, matrix, array, or list will expand that structure to accommodate the new value.**

```
x <- c(2, 6, 4)
```

```
x[7] <- 10
```

```
x
```

```
[1] 2 6 4 NA NA NA 10
```

The vector `x` has expanded from 3 elements to 7 elements through the assignment.

`x <- x[1:3]` would shrink it back to three elements again.

- **Indices in R start at 1, not at zero.**
In the vector above, `x[1]` is 2.



DATA MANAGEMENT

Topics

- arithmetic and logical operators
- creating, recoding, renaming variables
- missing values and data values
- type conversions
- sorting, merging, subsetting
- reshaping datasets

Arithmetic Operators

Operator

Description

+

Addition

-

Subtraction

*

Multiplication

/

Division

^ or **

Exponentiation

Logical Operators

Operator	Description
<	Less than
<=	Less than or equal to
>	Greater than
>=	Greater than or equal to
==	Exactly equal to
!=	Not equal to
!x	Not x
x y	x or y
x & y	x and y

Salaries dataset

```
data(Salaries, package="car")
```

```
names(Salaries)
```

```
[1] "rank"           "discipline"      "yrs.since.phd"  "yrs.service"  
[5] "sex"            "salary"
```

```
head(Salaries)
```

	rank	discipline	yrs.since.phd	yrs.service	sex	salary
1	Prof	B	19	18	Male	139750
2	Prof	B	20	16	Male	173200
3	AsstProf	B	4	3	Male	79750
4	Prof	B	45	39	Male	115000
5	Prof	B	40	41	Male	141500
6	AssocProf	B	6	6	Male	97000

Add new variables to a data frame

```
experience <- (yrs.since.phd + yrs.service)/2 # fails
```

```
experience <- (Salaries$yrs.since.phd +  
              Salaries$yrs.service)/2 # doesn't fail but...
```

```
Salaries$experience <- (Salaries$yrs.since.phd +  
                        Salaries$yrs.service)/2 # works
```

```
Salaries <- transform(Salaries, experience =  
                      (yrs.since.phd + yrs.service )/2) # better
```

Using the transform() function

```
Salaries <- transform(Salaries,  
  logSalary = log(salary),  
  experience = ( yrs.since.phd + yrs.service)/2,  
  discipline = factor(discipline, levels=c("A", "B"),  
    labels=c("Theoretical", "Applied")),  
  salaryCat = cut(salary,  
    quantile(salary, probs=c(0, .33, .66, 1)),  
    labels=c("low", "med", "high"))  
)
```

Renaming variables in a data frame

```
names(Salaries)
```

```
[1] "rank"          "discipline"    "yrs.since.phd" "yrs.service"  
[5] "sex"           "salary"
```

```
names(Salaries)[3:4] <- c("yrs.post.phd",  
                          "yrs.of.service")
```

```
names(Salaries)
```

```
[1] "rank"          "discipline"    "yrs.post.phd" "yrs.of.service"  
[5] "sex"           "salary"
```

Missing values

- coded as NA (no quotation marks)
- test with `is.na()`
- `x == NA` doesn't work

Working with missing values

```
data(sleep, package="VIM")
```

```
head(sleep)
```

```
head(is.na(sleep))
```

```
colSums(is.na(sleep))
```

```
colMeans(is.na(sleep))
```

```
newSleep <- na.omit(sleep)
```

Recoding to missing values

hypothetical example

- `df$age <- ifelse(df$age == 99, NA, df$age)`
- `df$age <- ifelse(df$age %in% c(99, 999, -1),
NA, df$age)`
- delete all missing values
`newdf <- na.omit(df) # listwise deletion`

Sorting data frames

```
index <- order(Salaries$rank, Salaries$salary)
Salaries <- Salaries[index, ]
```

- `order()` returns a permutation which rearranges its argument into ascending or descending order
- default is ascending (reverse with - sign)
- note the comma!

```
index <- order(Salaries$rank, -Salaries$salary)
newdf <- Salaries[index, ]
```

```
head(newdf)
```

	rank	discipline	yrs.since.phd	yrs.service	sex	salary
238	AsstProf	A	7	6	Female	63100
227	AsstProf	A	3	1	Male	63900
65	AsstProf	B	4	3	Male	68404
241	AsstProf	A	5	3	Male	69200
235	AsstProf	A	8	3	Male	69700
50	AsstProf	B	1	1	Male	70768

Merging data frames (horizontally)

Use merge()

```
dataframeC <- merge(dataframeA, dataframeB,  
  by = "ID")
```

```
dataframeC <- merge(dataframeA, dataframeB,  
  by=c("ID", "Country"))
```

Merging data frames (vertically)

Use `rbind()`

```
dataframeC <- rbind(dataframeA, dataframeB)
```

both data frames must have same variables (but don't have to be in same order).

Subsetting a data frame

Selecting (excluding variables)

```
df <- Salaries[c("rank", "sex", "salary")]
```

```
df <- Salaries[c(1, 5, 6)]
```

```
df <- df[-c(2, 3, 4)]
```

Subsetting a data frame

Selecting (excluding) observations

```
newdata <- Salaries[1:5, ]
```

```
newdata <- Salaries[Salaries$sex=="Female"  
                    & Salaries$salary > 100000, ]
```

Subsetting a data frame

Selecting observations/variables using subset()

```
newdata <- subset(Salaries,  
  salary >= 200000 | salary < 60000,  
  select=c(sex, rank, salary))
```

```
newdata <- subset(Salaries,  
  sex=="Female" & salary > 60000,  
  select=yrs.since.phd:salary)
```

Aggregating data

- `aggregate(x, by, FUN)`
 - *x* is the data object to be collapsed
 - *by* is a list of variables that will be cross to form new observations
 - *FUN* is a scalar function used to calculate summary statistics that will make up the new observation values

Aggregate example

```
aggdata <- aggregate(mtcars[c("mpg", "disp", "wt")],  
                     by=list(cylinder=mtcars$cyl, gears = mtcars$gear),  
                     FUN=mean, na.rm=TRUE)
```

aggdata

	cylinder	gears	mpg	disp	wt
1	4	3	21.5	120	2.46
2	6	3	19.8	242	3.34
3	8	3	15.1	358	4.10
4	4	4	26.9	103	2.38
5	6	4	19.8	164	3.09
6	4	5	28.2	108	1.83
7	6	5	19.7	145	2.77
8	8	5	15.4	326	3.37

Subgroup processing

- `by(data, INDICES, FUN)`
 - where *data* is a data frame or matrix
 - *INDICES* is a categorical variable or list of categorical variables that define the groups
 - *FUN* is an arbitrary function

Descriptive statistics by subgroup

```
by(mtcars[c("mpg", "hp")], mtcars$am, summary)
```

```
mtcars$am: 0
```

mpg	hp
Min. :10.4	Min. : 62
1st Qu.:14.9	1st Qu.:116
Median :17.3	Median :175
Mean :17.1	Mean :160
3rd Qu.:19.2	3rd Qu.:192
Max. :24.4	Max. :245

```
-----
```

```
mtcars$am: 1
```

mpg	hp
Min. :15	Min. : 52
1st Qu.:21	1st Qu.: 66
Median :23	Median :109
Mean :24	Mean :127
3rd Qu.:30	3rd Qu.:113
Max. :34	Max. :335

Reshaping wide to long

```
library(reshape2)
jtrain_long <- reshape(jtrain, idvar = "id",
                        varying = list(c("re74", "re75", "re78"),
                                       c("unem74", "unem75", "unem78")),
                        v.names = c("re", "unem"),
                        timevar = "year",
                        times = c(74, 75, 78),
                        direction = "long")
```

Reshaping long to wide

```
library(reshape2)
jtrain_wide<- reshape(jtrain_long,
                      idvar = "id",
                      v.names = c("re", "unem"),
                      timevar = "year",
                      direction = "wide")
```

Applying function to columns and rows

```
apply(dataframe, index, function, options)
```

```
apply(mtcars, 1, mean) # row means
```

```
apply(mtcars, 2, mean, na.rm=TRUE) # column means
```